

Lab 4 Report:

Surpass Human Performance in Fashion MNIST Classificaion

Name:

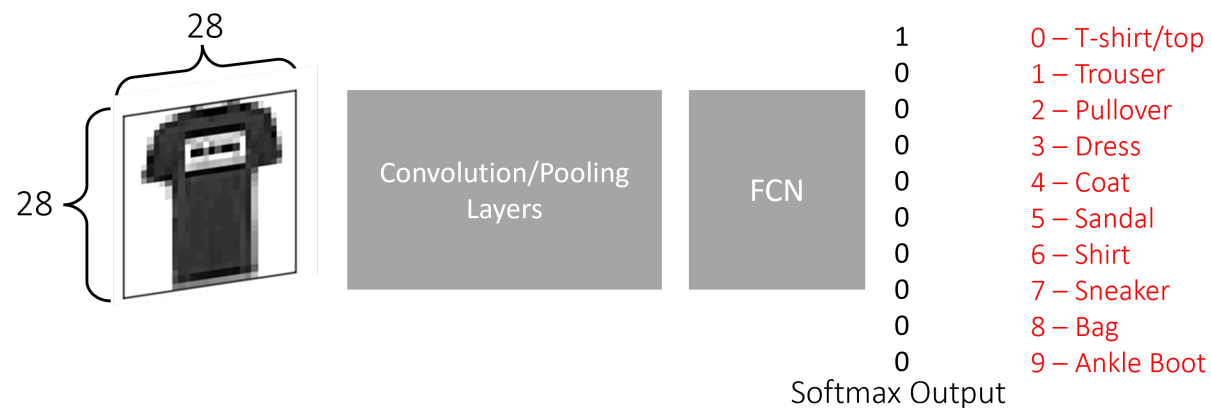
```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import torch
import numpy as np
```

```
In [2]: from IPython.display import Image # For displaying images in colab jupyter cell
```

```
In [3]: Image('lab4_exercise.png', width = 1000)
```

Out[3]:

Surpass Human Performance in Fashion MNIST Classification



In this exercise, you will classify fashion item class (28 x 28) using your own **Convolutional Neural Network Architecture**.

Prior to training your neural net, 1) Normalize the dataset using standard scaler and 2) Split the dataset into train/validation/test.

Design your own CNN architecture with your choice of Convolution/Pooling/FCN layers, activation functions, optimization method etc.

Your goal is to **achieve a testing accuracy of >89%**, with no restrictions on epochs (**Human performance: 83.5%**).

Demonstrate the performance of your model via plotting the **training loss**, **validation accuracy** and printing out the **testing accuracy**.

After your model has reached the goal, print the accuracy in each class. What is the class that your model performed the worst?

44

Prepare Data

```
In [4]: # Load Fashion-MNIST Dataset in Numpy

# 10000 training features/targets where each feature is a greyscale image with shape (28, 28)
train_features = np.load('fashion_mnist_train_features.npy')
train_targets = np.load('fashion_mnist_train_targets.npy')

# 1000 testing features/targets
```

```
test_features = np.load('fashion_mnist_test_features.npy')
test_targets = np.load('fashion_mnist_test_targets.npy')

# Let's see the shapes of training/testing datasets
print("Training Features Shape: ", train_features.shape)
print("Training Targets Shape: ", train_targets.shape)
print("Testing Features Shape: ", test_features.shape)
print("Testing Targets Shape: ", test_targets.shape)
```

Training Features Shape: (10000, 28, 28)

Training Targets Shape: (10000,)

Testing Features Shape: (1000, 28, 28)

Testing Targets Shape: (1000,)

In [5]: *# Visualizing the first three training features (samples)*


```
plt.figure(figsize = (10, 10))

plt.subplot(1,3,1)
plt.imshow(train_features[0], cmap = 'Greys')

plt.subplot(1,3,2)
plt.imshow(train_features[1], cmap = 'Greys')

plt.subplot(1,3,3)
plt.imshow(train_features[2], cmap = 'Greys')
```

Out[5]: <matplotlib.image.AxesImage at 0x21ae45bf520>



```
In [6]: # Reshape features via flattening the images
# This refers to reshape each sample from a 2d array to a 1d array.
# hint: np.reshape() function could be useful here
from sklearn.preprocessing import StandardScaler
import torch, torch.nn as nn

tr_flat = train_features.reshape(len(train_features), -1)
te_flat = test_features.reshape(len(test_features), -1)
scaler = StandardScaler()
```

```
In [7]: # Define your scaling function
# YOUR CODE HERE

# Scale the dataset according to standard scaling
train_features = scaler.fit_transform(tr_flat).reshape(-1,28,28)
test_features = scaler.transform(te_flat).reshape(-1,28,28)
```

```
In [8]: # Take the first 1000 (or randomly select 1000) training features and targets as validation set

val_features, val_targets = train_features[:1000], train_targets[:1000]

# Take the remaining 9000 training features and targets as training set
```

```
train_features, train_targets = train_features[1000:], train_targets[1000:]
```

In [9]: *# Reshape train/validation/test sets to conform to PyTorch's (N, Channels, Height, Width) standard for CNNs*

```
train_features = train_features[:, np.newaxis, :, :]      # (N,1,28,28)
val_features   = val_features[:, np.newaxis, :, :]      # (N_val,1,28,28)
test_features  = test_features[:, np.newaxis, :, :]     # (N_test,1,28,28)

print("Shapes after add channel axis:",
      train_features.shape, val_features.shape, test_features.shape)
```

Shapes after add channel axis: (9000, 1, 28, 28) (1000, 1, 28, 28) (1000, 1, 28, 28)

Define Model

In [10]: *# Define your CNN architecture here*

```
class CNNModel(nn.Module):
    """
    Two-conv modern CNN for Fashion-MNIST.
    Keeps the SAME spatial sizes as the course example (28→14→7) so the
    flattened feature vector is 32 × 7 × 7 = 1568.

    Layer order:
    Conv3×3 → BN → ReLU → MaxPool
    Conv3×3 → BN → ReLU → MaxPool
    Flatten → Dropout → FC-10 (logits)
    """
    def __init__(self):
        super().__init__()

        # ---- Block 1 -----
        # Input : (N, 1, 28, 28)
        # Output: (N, 16, 28, 28) - same H/W thanks to padding=1
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=16,
                               kernel_size=3, padding=1)
        self.bn1   = nn.BatchNorm2d(16)           # stabilise activations
        self.pool1 = nn.MaxPool2d(kernel_size=2) # 28 → 14
```

```

# ---- Block 2 -----
# Input : (N, 16, 14, 14)
# Output: (N, 32, 14, 14)
self.conv2 = nn.Conv2d(in_channels=16, out_channels=32,
                       kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(32)
self.pool2 = nn.MaxPool2d(kernel_size=2) # 14 → 7

# ---- Classifier -----
# Flatten: (N, 32, 7, 7) → (N, 1568)
self.dropout = nn.Dropout(p=0.4) # reduce over-fitting
self.fc1 = nn.Linear(32 * 7 * 7, 10) # 10 logits for CE Loss

def forward(self, x):
    """
    Forward pass:
    x : Tensor of shape (batch, 1, 28, 28)
    returns raw logits (no soft-max; CrossEntropyLoss handles it)
    """
    # Block 1
    x = self.conv1(x) # convolution
    x = self.bn1(x) # batch normalisation
    x = torch.relu(x) # non-linearity
    x = self.pool1(x) # spatial down-sampling → 14×14

    # Block 2
    x = self.conv2(x)
    x = self.bn2(x)
    x = torch.relu(x)
    x = self.pool2(x) # → 7×7

    # Classifier
    x = x.view(x.size(0), -1) # flatten to (N, 1568)
    x = self.dropout(x)
    logits = self.fc1(x)

    return logits

```

Select Hyperparameters

In [11]: *# Fix the random seed so that model performance is reproducible*

```

DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'
print("Using device:", DEVICE)
torch.manual_seed(55)

# Initialize your CNN model

model = CNNModel().to(DEVICE)

# Define learning rate, epoch and batchsize for mini-batch gradient

learning_rate = 1e-4      # Adam → 1 × 10-4
epochs        = 250
batchsize     = 128

# Define loss function and optimizer

loss_func = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

model

```

Using device: cuda

Out[11]: CNNModel(
 (conv1): Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
 (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
 (dropout): Dropout(p=0.4, inplace=False)
 (fc1): Linear(in_features=1568, out_features=10, bias=True)
)

Identify Tracked Values

```
In [12]: # Placeholders for training loss and validation accuracy during training
# Training loss should be tracked for each iteration (1 iteration -> single forward pass to the network)
# Validation accuracy should be evaluated every 'Epoch' (1 epoch -> full training dataset)
# If using batch gradient, 1 iteration = 1 epoch

train_loss_list      = []
val_accuracy_list    = np.zeros(epochs)
```

Train Model

```
In [13]: import tqdm # Use "for epoch in tqdm.trange(epochs):" to see the progress bar

# Convert the training, validation, testing dataset (NumPy arrays) into torch tensors
# Split your training features/targets into mini-batches if using mini-batch gradient

# -----
# Convert NumPy → Torch tensors, move to DEVICE (CPU or GPU)
# -----
tr_x = torch.from_numpy(train_features).float().to(DEVICE) # training images
tr_y = torch.from_numpy(train_targets).long().to(DEVICE)  # training labels
val_x = torch.from_numpy(val_features).float().to(DEVICE) # validation images
val_y = torch.from_numpy(val_targets).long().to(DEVICE)   # validation labels
te_x  = torch.from_numpy(test_features).float().to(DEVICE) # test images
te_y  = torch.from_numpy(test_targets).long().to(DEVICE)   # test labels

# Split the training set into mini-batches (128 samples each here)
tr_batches_x = torch.split(tr_x, batchsize)
tr_batches_y = torch.split(tr_y, batchsize)

# -----
# Training Loop
# -----
print("Model on:", next(model.parameters()).device) # sanity-check: should print cuda:0

for epoch in tqdm.trange(epochs, desc='Epoch'):
```



```

# ----- train mode -----
model.train()
for xb, yb in zip(tr_batches_x, tr_batches_y):
    optimizer.zero_grad()          # reset accumulated gradients
    logits = model(xb)             # forward pass
    loss = loss_func(logits, yb)   # compute cross-entropy
    train_loss_list.append(loss.item()) # stash for plotting
    loss.backward()                # back-propagate
    optimizer.step()               # update weights

# ----- validation -----
model.eval()                      # turn off dropout / BN updates
with torch.no_grad():             # no grad tracking ⇒ Lower memory
    preds = model(val_x).argmax(dim=1) # get predicted class indices
    acc = (preds == val_y).float().mean().item()
    val_accuracy_list[epoch] = acc

# Console feedback per epoch
print(f"Epoch {epoch:02d} Val Acc = {acc*100:5.2f}%")

```

Model on: cuda:0

Epoch: 1%| | 2/250 [00:00<00:46, 5.37it/s]

Epoch 00 Val Acc = 72.10%

Epoch 01 Val Acc = 75.30%

Epoch: 1%| | 3/250 [00:00<00:36, 6.76it/s]

Epoch 02 Val Acc = 77.80%

Epoch 03 Val Acc = 79.70%

Epoch: 2%|| | 6/250 [00:00<00:29, 8.25it/s]

Epoch 04 Val Acc = 80.80%

Epoch 05 Val Acc = 82.20%

Epoch: 4%|| | 9/250 [00:01<00:25, 9.46it/s]

Epoch 06 Val Acc = 82.90%

Epoch 07 Val Acc = 83.50%

Epoch 08 Val Acc = 84.50%

Epoch: 5%|| | 12/250 [00:01<00:22, 10.47it/s]

Epoch 09 Val Acc = 84.70%


Epoch 10 Val Acc = 85.10%

Epoch 11 Val Acc = 85.30%

Epoch: 6%|| | 14/250 [00:01<00:22, 10.36it/s]

Epoch 12 Val Acc = 85.70%


Epoch 13 Val Acc = 85.60%

Epoch: 6% |  | 16/250 [00:01<00:23, 10.17it/s]

Epoch 14 Val Acc = 86.10%

Epoch 15 Val Acc = 86.60%


Epoch 16 Val Acc = 86.60%

Epoch: 8% |  | 20/250 [00:02<00:21, 10.86it/s]

Epoch 17 Val Acc = 86.80%

Epoch 18 Val Acc = 87.00%


Epoch 19 Val Acc = 87.00%

Epoch: 9% |  | 22/250 [00:02<00:20, 11.25it/s]

Epoch 20 Val Acc = 87.20%


Epoch 21 Val Acc = 87.10%

Epoch 22 Val Acc = 87.50%

Epoch: 10% |  | 24/250 [00:02<00:19, 11.72it/s]


Epoch 23 Val Acc = 87.30%

Epoch 24 Val Acc = 87.10%

Epoch: 10% |  | 26/250 [00:02<00:19, 11.21it/s]

Epoch 25 Val Acc = 87.30%


Epoch 26 Val Acc = 87.70%

Epoch: 12% |  | 30/250 [00:03<00:20, 10.61it/s]

Epoch 27 Val Acc = 87.60%

Epoch 28 Val Acc = 87.90%


Epoch 29 Val Acc = 88.50%

Epoch: 13% |  | 32/250 [00:03<00:19, 10.97it/s]

Epoch 30 Val Acc = 88.20%

Epoch 31 Val Acc = 88.40%


Epoch 32 Val Acc = 88.50%

Epoch: 14% |  | 36/250 [00:03<00:19, 11.22it/s]

Epoch 33 Val Acc = 88.30%

Epoch 34 Val Acc = 88.50%


Epoch 35 Val Acc = 88.90%

Epoch: 15% |  | 38/250 [00:03<00:19, 10.74it/s]


Epoch 36 Val Acc = 88.60%

Epoch 37 Val Acc = 88.70%


Epoch 38 Val Acc = 89.00%

Epoch: 17% |  | 42/250 [00:04<00:19, 10.93it/s]


Epoch 39 Val Acc = 89.10%
Epoch 40 Val Acc = 88.90%
Epoch 41 Val Acc = 88.80%

Epoch: 18% |  | 44/250 [00:04<00:18, 11.28it/s]


Epoch 42 Val Acc = 89.00%
Epoch 43 Val Acc = 88.70%
Epoch 44 Val Acc = 89.00%

Epoch: 19% |  | 48/250 [00:04<00:17, 11.41it/s]


Epoch 45 Val Acc = 89.10%
Epoch 46 Val Acc = 88.80%
Epoch 47 Val Acc = 89.20%

Epoch: 20% |  | 50/250 [00:04<00:19, 10.29it/s]


Epoch 48 Val Acc = 89.20%
Epoch 49 Val Acc = 89.00%

Epoch: 21% |  | 52/250 [00:05<00:18, 10.45it/s]


Epoch 50 Val Acc = 89.20%
Epoch 51 Val Acc = 88.80%
Epoch 52 Val Acc = 89.00%

Epoch: 22% |  | 56/250 [00:05<00:17, 11.27it/s]


Epoch 53 Val Acc = 88.80%
Epoch 54 Val Acc = 88.80%
Epoch 55 Val Acc = 89.30%

Epoch: 23% |  | 58/250 [00:05<00:16, 11.56it/s]


Epoch 56 Val Acc = 89.10%
Epoch 57 Val Acc = 89.00%
Epoch 58 Val Acc = 88.90%

Epoch: 25% |  | 62/250 [00:05<00:17, 10.69it/s]


Epoch 59 Val Acc = 88.90%
Epoch 60 Val Acc = 89.20%
Epoch 61 Val Acc = 89.20%

Epoch: 26% |  | 64/250 [00:06<00:18, 10.30it/s]

Epoch 62 Val Acc = 89.10%
Epoch 63 Val Acc = 89.00%

Epoch: 26% |  | 66/250 [00:06<00:18, 10.03it/s]

Epoch 64 Val Acc = 89.20%
Epoch 65 Val Acc = 89.60%
Epoch 66 Val Acc = 89.40%

Epoch: 27% |  | 68/250 [00:06<00:18, 9.93it/s]

Epoch 67 Val Acc = 89.40%

Epoch 68 Val Acc = 89.20%

Epoch: 28% |██████| 71/250 [00:06<00:18, 9.54it/s]

Epoch 69 Val Acc = 89.40%

Epoch 70 Val Acc = 89.20%

Epoch 71 Val Acc = 89.20%

Epoch: 30% |██████| 75/250 [00:07<00:16, 10.89it/s]

Epoch 72 Val Acc = 89.30%

Epoch 73 Val Acc = 89.90%

Epoch 74 Val Acc = 89.60%

Epoch: 31% |██████| 77/250 [00:07<00:16, 10.66it/s]

Epoch 75 Val Acc = 89.60%

Epoch 76 Val Acc = 89.50%

Epoch: 32% |██████| 79/250 [00:07<00:16, 10.27it/s]

Epoch 77 Val Acc = 89.50%

Epoch 78 Val Acc = 89.50%

Epoch: 32% |██████| 81/250 [00:07<00:16, 10.00it/s]

Epoch 79 Val Acc = 89.70%

Epoch 80 Val Acc = 89.50%

Epoch: 33% |██████| 83/250 [00:08<00:16, 9.95it/s]

Epoch 81 Val Acc = 89.40%

Epoch 82 Val Acc = 89.60%

Epoch: 34% |██████| 85/250 [00:08<00:17, 9.70it/s]

Epoch 83 Val Acc = 89.50%

Epoch 84 Val Acc = 89.50%

Epoch 85 Val Acc = 89.60%

Epoch: 36% |██████| 89/250 [00:08<00:15, 10.35it/s]

Epoch 86 Val Acc = 89.80%

Epoch 87 Val Acc = 89.60%

Epoch 88 Val Acc = 89.80%

Epoch: 36% |██████| 91/250 [00:08<00:15, 10.05it/s]

Epoch 89 Val Acc = 89.70%

Epoch 90 Val Acc = 89.40%

Epoch: 37% |██████| 93/250 [00:09<00:15, 10.16it/s]

Epoch 91 Val Acc = 89.90%

Epoch 92 Val Acc = 89.80%

Epoch 93 Val Acc = 89.80%

Epoch: 39% |██████| 97/250 [00:09<00:14, 10.91it/s]

Epoch 94 Val Acc = 89.80%

Epoch 95 Val Acc = 89.50%

Epoch 96 Val Acc = 89.70%

Epoch: 40% |██████| 99/250 [00:09<00:14, 10.36it/s]

Epoch 97 Val Acc = 89.70%

Epoch 98 Val Acc = 89.50%

Epoch: 40% |██████| 101/250 [00:09<00:14, 10.17it/s]

Epoch 99 Val Acc = 89.60%

Epoch 100 Val Acc = 90.20%

Epoch: 41% |██████| 103/250 [00:10<00:14, 10.35it/s]

Epoch 101 Val Acc = 90.10%

Epoch 102 Val Acc = 89.80%

Epoch 103 Val Acc = 89.90%

Epoch: 43% |██████| 107/250 [00:10<00:13, 10.67it/s]

Epoch 104 Val Acc = 90.10%

Epoch 105 Val Acc = 89.60%

Epoch 106 Val Acc = 90.00%

Epoch: 44% |██████| 109/250 [00:10<00:13, 10.36it/s]

Epoch 107 Val Acc = 90.00%

Epoch 108 Val Acc = 90.10%

Epoch: 44% |██████| 111/250 [00:10<00:14, 9.87it/s]

Epoch 109 Val Acc = 89.90%

Epoch 110 Val Acc = 89.70%

Epoch: 46% |██████| 114/250 [00:11<00:13, 10.28it/s]

Epoch 111 Val Acc = 89.50%

Epoch 112 Val Acc = 89.90%

Epoch 113 Val Acc = 90.10%

Epoch: 46% |██████| 116/250 [00:11<00:12, 10.83it/s]

Epoch 114 Val Acc = 89.80%

Epoch 115 Val Acc = 89.60%

Epoch 116 Val Acc = 90.00%

Epoch: 48% |██████| 120/250 [00:11<00:11, 11.53it/s]

Epoch 117 Val Acc = 89.80%

Epoch 118 Val Acc = 89.80%

Epoch 119 Val Acc = 89.40%

Epoch: 49% |██████| 122/250 [00:11<00:12, 10.65it/s]

```
Epoch 120 Val Acc = 89.70%
Epoch 121 Val Acc = 89.90%
Epoch: 50%|██████| 124/250 [00:12<00:12, 10.12it/s]
Epoch 122 Val Acc = 90.20%
Epoch 123 Val Acc = 90.20%
Epoch: 50%|██████| 126/250 [00:12<00:12, 9.99it/s]
Epoch 124 Val Acc = 89.70%
Epoch 125 Val Acc = 89.70%
Epoch: 51%|██████| 128/250 [00:12<00:12, 9.93it/s]
Epoch 126 Val Acc = 89.90%
Epoch 127 Val Acc = 89.60%
Epoch 128 Val Acc = 89.40%
Epoch: 52%|██████| 131/250 [00:12<00:12, 9.52it/s]
Epoch 129 Val Acc = 90.00%
Epoch 130 Val Acc = 90.00%
Epoch: 53%|██████| 133/250 [00:12<00:12, 9.58it/s]
Epoch 131 Val Acc = 90.00%
Epoch 132 Val Acc = 89.90%
Epoch 133 Val Acc = 89.90%
Epoch: 55%|██████| 137/250 [00:13<00:10, 10.56it/s]
Epoch 134 Val Acc = 89.90%
Epoch 135 Val Acc = 90.00%
Epoch 136 Val Acc = 90.00%
Epoch: 56%|██████| 139/250 [00:13<00:10, 10.46it/s]
Epoch 137 Val Acc = 89.70%
Epoch 138 Val Acc = 90.10%
Epoch 139 Val Acc = 89.60%
Epoch: 56%|██████| 141/250 [00:13<00:09, 10.97it/s]
Epoch 140 Val Acc = 90.00%
Epoch 141 Val Acc = 89.80%
Epoch: 58%|██████| 145/250 [00:14<00:10, 10.47it/s]
Epoch 142 Val Acc = 90.10%
Epoch 143 Val Acc = 90.10%
Epoch 144 Val Acc = 89.80%
Epoch: 59%|██████| 147/250 [00:14<00:09, 10.56it/s]
Epoch 145 Val Acc = 90.10%
Epoch 146 Val Acc = 90.00%
Epoch 147 Val Acc = 89.90%
```

Epoch: 60% |███████| | 151/250 [00:14<00:09, 10.37it/s]

Epoch 148 Val Acc = 90.10%

Epoch 149 Val Acc = 90.00%

Epoch 150 Val Acc = 90.00%

Epoch: 61% |███████| | 153/250 [00:14<00:09, 9.74it/s]

Epoch 151 Val Acc = 89.70%

Epoch 152 Val Acc = 90.20%

Epoch: 62% |███████| | 156/250 [00:15<00:09, 10.38it/s]

Epoch 153 Val Acc = 90.10%

Epoch 154 Val Acc = 90.50%

Epoch 155 Val Acc = 90.60%

Epoch: 63% |███████| | 158/250 [00:15<00:08, 10.97it/s]

Epoch 156 Val Acc = 90.10%

Epoch 157 Val Acc = 89.90%

Epoch 158 Val Acc = 90.30%

Epoch: 64% |███████| | 160/250 [00:15<00:08, 10.71it/s]

Epoch 159 Val Acc = 90.40%

Epoch 160 Val Acc = 90.10%

Epoch: 66% |███████| | 164/250 [00:15<00:08, 10.22it/s]

Epoch 161 Val Acc = 90.10%

Epoch 162 Val Acc = 90.40%

Epoch 163 Val Acc = 90.00%

Epoch: 66% |███████| | 166/250 [00:16<00:07, 10.54it/s]

Epoch 164 Val Acc = 89.90%

Epoch 165 Val Acc = 90.70%

Epoch 166 Val Acc = 90.20%

Epoch: 67% |███████| | 168/250 [00:16<00:07, 10.61it/s]

Epoch 167 Val Acc = 90.30%

Epoch 168 Val Acc = 90.60%

Epoch: 69% |███████| | 172/250 [00:16<00:07, 10.98it/s]

Epoch 169 Val Acc = 90.00%

Epoch 170 Val Acc = 89.80%

Epoch 171 Val Acc = 90.00%

Epoch: 70% |███████| | 174/250 [00:16<00:07, 10.36it/s]

Epoch 172 Val Acc = 90.10%

Epoch 173 Val Acc = 90.40%

Epoch: 70% |███████| | 176/250 [00:17<00:07, 10.37it/s]

```
Epoch 174 Val Acc = 90.10%
Epoch 175 Val Acc = 89.90%
Epoch 176 Val Acc = 90.40%
Epoch: 72%|███████| 180/250 [00:17<00:06, 11.29it/s]
Epoch 177 Val Acc = 90.40%
Epoch 178 Val Acc = 90.10%
Epoch 179 Val Acc = 89.90%
Epoch: 73%|███████| 182/250 [00:17<00:05, 11.54it/s]
Epoch 180 Val Acc = 90.20%
Epoch 181 Val Acc = 90.00%
Epoch 182 Val Acc = 90.30%
Epoch: 74%|███████| 186/250 [00:17<00:05, 11.02it/s]
Epoch 183 Val Acc = 90.10%
Epoch 184 Val Acc = 90.10%
Epoch 185 Val Acc = 90.40%
Epoch: 75%|███████| 188/250 [00:18<00:05, 11.14it/s]
Epoch 186 Val Acc = 89.90%
Epoch 187 Val Acc = 90.20%
Epoch 188 Val Acc = 90.10%
Epoch: 77%|███████| 192/250 [00:18<00:04, 11.81it/s]
Epoch 189 Val Acc = 90.00%
Epoch 190 Val Acc = 89.90%
Epoch 191 Val Acc = 90.30%
Epoch: 78%|███████| 194/250 [00:18<00:04, 12.06it/s]
Epoch 192 Val Acc = 90.00%
Epoch 193 Val Acc = 90.00%
Epoch 194 Val Acc = 90.20%
Epoch: 78%|███████| 196/250 [00:18<00:04, 11.85it/s]
Epoch 195 Val Acc = 90.00%
Epoch 196 Val Acc = 90.20%
Epoch: 80%|███████| 200/250 [00:19<00:04, 11.25it/s]
Epoch 197 Val Acc = 89.90%
Epoch 198 Val Acc = 89.70%
Epoch 199 Val Acc = 89.50%
Epoch: 81%|███████| 202/250 [00:19<00:04, 11.65it/s]
Epoch 200 Val Acc = 90.00%
Epoch 201 Val Acc = 90.20%
Epoch 202 Val Acc = 89.90%
```


Epoch: 82% |██████████| | 204/250 [00:19<00:04, 11.20it/s]

Epoch 203 Val Acc = 89.70%

Epoch 204 Val Acc = 90.10%

Epoch: 82% |██████████| | 206/250 [00:19<00:04, 10.60it/s]

Epoch 205 Val Acc = 90.00%

Epoch 206 Val Acc = 90.10%

Epoch: 84% |██████████| | 210/250 [00:20<00:03, 10.71it/s]

Epoch 207 Val Acc = 90.20%

Epoch 208 Val Acc = 90.00%

Epoch 209 Val Acc = 90.10%

Epoch: 85% |██████████| | 212/250 [00:20<00:03, 10.95it/s]

Epoch 210 Val Acc = 89.90%

Epoch 211 Val Acc = 90.10%

Epoch 212 Val Acc = 89.90%

Epoch: 86% |██████████| | 216/250 [00:20<00:03, 11.16it/s]

Epoch 213 Val Acc = 89.90%

Epoch 214 Val Acc = 90.20%

Epoch 215 Val Acc = 89.90%

Epoch: 87% |██████████| | 218/250 [00:20<00:02, 10.86it/s]

Epoch 216 Val Acc = 90.00%

Epoch 217 Val Acc = 90.10%

Epoch: 88% |██████████| | 220/250 [00:20<00:02, 10.37it/s]

Epoch 218 Val Acc = 89.90%

Epoch 219 Val Acc = 89.80%

Epoch 220 Val Acc = 89.70%

Epoch: 90% |██████████| | 224/250 [00:21<00:02, 11.23it/s]

Epoch 221 Val Acc = 90.00%

Epoch 222 Val Acc = 90.10%

Epoch 223 Val Acc = 89.70%

Epoch: 90% |██████████| | 226/250 [00:21<00:02, 11.63it/s]

Epoch 224 Val Acc = 90.00%

Epoch 225 Val Acc = 90.00%

Epoch 226 Val Acc = 89.50%

Epoch: 92% |██████████| | 230/250 [00:21<00:01, 11.95it/s]

Epoch 227 Val Acc = 90.30%

Epoch 228 Val Acc = 89.70%

Epoch 229 Val Acc = 90.10%

```
Epoch: 93%|██████████| 232/250 [00:21<00:01, 12.22it/s]
Epoch 230 Val Acc = 90.20%
Epoch 231 Val Acc = 89.90%
Epoch 232 Val Acc = 90.50%
Epoch: 94%|██████████| 236/250 [00:22<00:01, 12.16it/s]
Epoch 233 Val Acc = 89.90%
Epoch 234 Val Acc = 90.10%
Epoch 235 Val Acc = 90.00%
Epoch: 95%|██████████| 238/250 [00:22<00:00, 12.18it/s]
Epoch 236 Val Acc = 89.90%
Epoch 237 Val Acc = 90.40%
Epoch 238 Val Acc = 90.40%
Epoch: 97%|██████████| 242/250 [00:22<00:00, 11.52it/s]
Epoch 239 Val Acc = 90.30%
Epoch 240 Val Acc = 90.30%
Epoch 241 Val Acc = 89.80%
Epoch: 98%|██████████| 244/250 [00:23<00:00, 10.94it/s]
Epoch 242 Val Acc = 89.90%
Epoch 243 Val Acc = 89.90%
Epoch 244 Val Acc = 89.90%
Epoch: 99%|██████████| 248/250 [00:23<00:00, 11.75it/s]
Epoch 245 Val Acc = 90.20%
Epoch 246 Val Acc = 90.10%
Epoch 247 Val Acc = 90.30%
Epoch: 100%|██████████| 250/250 [00:23<00:00, 10.64it/s]
Epoch 248 Val Acc = 90.20%
Epoch 249 Val Acc = 90.00%
```

Visualize & Evaluate Model

```
In [14]: # Seaborn for prettier plot
```

```
import seaborn as sns
```


```
sns.set(style = 'whitegrid', font_scale = 1)
```

```
In [15]: # Visualize training loss

plt.figure(figsize = (15, 9))

plt.subplot(2, 1, 1)
plt.plot(train_loss_list, linewidth = 3)
plt.ylabel("training loss")
plt.xlabel("iterations")
sns.despine()

plt.subplot(2, 1, 2)
plt.plot(val_accuracy_list, linewidth = 3, color = 'gold')
plt.ylabel("validation accuracy")
sns.despine()
```



We found that the learning curve when learning rate is e^{-3} is spiking, which may be due to the poor choice of learning rate.

After changing it to e^{-4} , the curve is smoother

In [16]: *# Compute the testing accuracy*

```
# YOUR CODE HERE
# -----
#   Evaluate the final model on the *held-out* test set
# -----
model.eval()                                # inference mode: no dropout/B-N updates
with torch.no_grad():                       # disable autograd → lower memory
    logits = model(te_x)                   # forward pass on all test images
    test_preds = logits.argmax(dim=1)       # predicted class indices
    test_acc = (test_preds == te_y).float().mean().item()

print(f"\nOverall Test Accuracy: {test_acc*100:.2f}%")

# -----
#   Per-class accuracy analysis
# -----
labels = [
    'T-shirt', 'Trouser', 'Pullover', 'Dress', 'Coat',
    'Sandal', 'Shirt', 'Sneaker', 'Bag', 'AnkleBoot'
]

cls_corr = [0] * 10                        # correct predictions per class
cls_tot = [0] * 10                         # total samples per class

# Tally results class-by-class
for true_lbl, pred_lbl in zip(te_y, test_preds):
    cls_tot[true_lbl] += 1
    cls_corr[true_lbl] += int(pred_lbl == true_lbl)
```

Overall Test Accuracy: 89.20%

In [17]: *# (OPTIONAL) Print the testing accuracy for each fashion class. Your code should produce something that looks like:*
Clever usage of np.where() could be useful here

```
# "Accuracy of T-shirt/top: 93.5 %"
# "Accuracy of Trouser: 89.3 %"
# etc...

# What's the fashion item that your model had the hardest time classifying?
```

```
accuracies = [c/t for c, t in zip(cls_corr, cls_tot)]  
worst = int(np.argmax(accuracies))  
# YOUR CODE HERE  
for i, lbl in enumerate(labels):  
    print(f"{lbl:10s}: {cls_corr[i]/cls_tot[i]*100:6.2f}%")  
print(f"\nWorst-classified item: {labels[worst]}")
```

T-shirt : 88.79%
Trouser : 99.05%
Pullover : 85.59%
Dress : 83.87%
Coat : 80.87%
Sandal : 95.40%
Shirt : 73.20%
Sneaker : 95.79%
Bag : 96.84%
AnkleBoot : 94.74%

Worst-classified item: Shirt